INDEPENDENT CLOUD-READINESS REVIEW

# Evidence before the estimate.

PreFlight is an offline, evidence-first review of a Guidewire-style policy administration estate. It inventories what exists, links claims to files, records every disposition, and names every unknown — before a migration estimate is signed.

OFFLINE · SYNTHETIC ESTATE · INDEPENDENT RESEARCH PROTOTYPE · NOT AFFILIATED WITH GUIDEWIRE

v1.0.0 · run R-DEMO-001 · demo-slice-300 · 2026-11-09

**REVIEW PACKET**
PKT-R-DEMO-001

6 SECTIONS

| | |
|---|---|
| Files hashed and classified | 300 |
| Features traced | 8 |
| Evidence records | 126 |
| Advisory findings / rules | 24 / 8 |
| Explicit dispositions | 12 |
| Audit events | 9 seed · 11 after 2 actions |

THE PROBLEM

ESTIMATE-MOVING SCOPE

# A migration estimate is a promise. The inventory is the proof.

The vendor estimate is not the risk by itself. The risk is the scope hidden underneath it: files that were never classified, features nobody traced, fields nobody disposed, and questions nobody owns.

> **Six to eight weeks** is a vendor estimate for APD product conversion — not a full migration estimate. It is a useful anchor. It is not an inventory of *your* estate.

**01**
### What exists?
Every file, hash, parser status, and the four that cannot be parsed — listed, not dropped.

**02**
### Which feature?
Cross-artefact links from PCF to entity, typelist, template, Gosu, and tests.

**03**
### What is disposed?
An explicit source-to-target disposition, or a named gap. Never a silent rename.

**04**
### What is unknown?
Owner, impact range, and the one question that would close the gap.

APD TYPES

### 24-character type limit → compile errors
APD limits **type names** to 24 characters. A longer name can be valid in source and still fail when endpoint generation compiles the generated code. The documented correction is a codegen `nameOverride` with the installed name.

**Guidewire docs** · “Correcting compile errors” · accessed 2026-11-09

PATTERNS

### JDBC, explicit commits, servlets, hidden packages
These recur in migration lessons. They often surface in the CUE (the official Cloud Upgrade Estimate) or a build gate; the earlier they are found, the less the estimate moves.

**Guidewire developer blog** · migration lessons guest post · accessed 2026-11-09

SBT

### Standards Based Template lines cannot convert
The Olos release note says SBT lines cannot be converted to APD-managed lines “at this time.” That is a readiness **scope and estimate risk**: it has to be priced or explicitly excluded, not assumed away.

**Guidewire release highlights** · Olos · accessed 2026-11-09

HANDOVER

### Incomplete handover
Inventory lives in spreadsheets, heads, and ticket threads. Dispositions and assumptions are rarely written with an owner and a falsifier.

**Field observation** · not measured in this study

**01**
**Scope becomes countable**
300 files, 4 exclusions, 8 features, 126 evidence records.

**02**
**Dispositions become explicit**
12 source-to-target rows, 5 target changes, 0 `_gen` writes.

**03**
**Unknowns become questions**
Each gap gets an owner, an impact range, and a question to ask.

Source discipline: vendor claims are labelled as vendor claims; gated bodies are marked gated; “incomplete handover” is a design premise, not a measured result. Full list in `docs/presentation/sources.md`.

THE GAP

CAPABILITY MAP, NOT A BENCHMARK

# No public tool we found combines PCF/entity/typelist/template/Gosu evidence into one disposition + assumption packet.

Guidewire CUE, Cloud Assurance Services, Lighthouse, SurePath, Developer Assistants, GoQu, CAST/ADDI for portfolio and code discovery, Syntra ETL for data migration and extraction, and SI assessments each cover a real slice. The gap is the hand-off: a file-level, cross-artefact feature inventory with a disposition ledger, an assumption register, and a reviewable packet.

CUE · CLOUD ASSURANCE SERVICES · LIGHTHOUSE · SUREPATH · DEVELOPER ASSISTANTS · GOQU · CAST / ADDI · SYNTRA ETL · SI ASSESSMENTS

| TOOL / SERVICE | INVENTORY | FEATURE EVIDENCE | DISPOSITION LEDGER | ASSUMPTION REGISTER | REVIEW PACKET |
|---|---|---|---|---|---|
| **Cloud Assurance Services**<br>official standards gate | code scan | no | no | no | findings |
| **Guidewire CUE**<br>official Cloud Upgrade Estimate | scope | no | no | estimate risks | estimate |
| **Lighthouse / SurePath / Assistants**<br>migration lab, delivery method, codegen help | lab | validated flows | no | no | handover |
| **SonarQube / GoQu**<br>Gosu lint and quality | code only | no | no | no | issues |
| **CAST / ADDI**<br>portfolio and code discovery | portfolio / code | dependencies | no | no | no |
| **Syntra ETL**<br>data migration / extraction | data | no | no | no | no |
| **Systems integrator**<br>services and expertise | bespoke | varies | varies | varies | varies |
| **PreFlight**<br>offline, one-feature evidence review | 300 files | cross-artefact | 12 explicit | impact-rated | 6 sections |

THE ACTUAL GAP

### The connective tissue, not another scanner.

Cross-artefact feature evidence + an explicit disposition ledger + an assumption register + a review packet. Each is adjacent to an existing tool; the integrated, reviewable chain is the missing artifact. Output: feature map · disposition ledger · assumption register · review packet.

DELIBERATELY NOT REPLACED

PreFlight does not compete with the official gate, the official estimate, the migration lab, or SI delivery.

CLOUD ASSURANCE · CUE · LIGHTHOUSE · SUREPATH · APD CONVERSION · CLOUD RETROFIT · SI DELIVERY

**Positioning, not an independent product evaluation.** Public product descriptions and documentation; see `sources.md` S-05, S-06, S-07, S-19–S-24.

THE ENGINEERED SOLUTION

FASTAPI · PYTHON 3.12 · JSON STORE

# One pipeline. Nine stages. Every step leaves a trace.

A read-only snapshot is hashed, parsed, assembled into features, reviewed against cited rules, reconstructed into a local draft, compared, evaluated, and sealed into a review packet.

| 01 | 02 | 03 | 04 | 05 | 06 | 07 | 08 | 09 |
|---|---|---|---|---|---|---|---|---|
| **Snapshot** | **Intake + hash** | **Parse** | **Discover** | **Review rules** | **Reconstruct** | **Compare** | **Evaluate** | **Packet** |
| read-only tree; no live tenant | manifest, status, before/after hash | PCF, entity, typelist, template, Gosu | features, evidence, gaps, confidence | 8 rules, 24 hard negatives | _ext draft, mapping, assumptions | preserved, changed, missing, unresolved | M1–M6, baselines, intervals | 6 sections, hashes, audit head |

**Read-only**
Snapshot hashed before and after; source unchanged.

**Unsupported listed**
Binary, corrupt, and unsupported files are named.

**Every claim evidence**
Path, locator, excerpt, hash, method, confidence.

**Every change pending**
Draft changes wait for an explicit reviewer action.

**No LLM required**
Model off; deterministic retrieval; no external calls.

**300** files → **126** evidence records → **24** findings / 8 rules → **12** dispositions → **1** review packet

one feature slice · one laptop · no telemetry

**INPUT**
**Read-only snapshot**
300 synthetic files; 4 excluded and named; source hash identical before and after.

**OUTPUT**
**Evidence graph + packet**
126 evidence records, 8 features, 24 cited findings, 12 dispositions, 1 packet.

**BOUNDARY**
**Stops before deployment**
No tenant, no LLM, no external calls, no `_gen` edit, no certification.

## WHAT EACH STAGE WRITES

| | | | | | |
|---|---|---|---|---|---|
| **SNAPSHOT** | read-only source; no writes | **INTAKE** | manifest · immutability · audit events | **PARSE** | evidence records with locators |
| **DISCOVER** | features · gaps · spine | **REVIEW** | findings · review evidence | **RECONSTRUCT** | mappings · changes · assumptions · draft |
| **COMPARE** | comparison rows | **EVALUATE** | metrics · intervals · baselines | **PACKET** | packet JSON/HTML · audit head |

**DETERMINISTIC** **Same snapshot + same rule pack → same evidence.** Reviewer decisions are the only mutation; every decision is hashed into the audit chain.

# A claim is only as strong as its locator, its label, and its open question.

PreFlight does not store “a finding.” It stores a claim with a source pointer, an excerpt hash, a method, a confidence label, and — where the evidence runs out — an assumption with an owner.

## ANATOMY OF A CLAIM

```
{ "evidence_id": "EV-0001", "path":
"config/web/pcf/policy/PAVehicleDV.pcf", "locator": "xpath
/PCF/Screen/TypeKeyInput[1]:4", "excerpt": "<TypeKeyInput
id=\"VehicleType\" …/>", "excerpt_sha": "d76bf9908b1f…",
"method": "xml_xpath", "confidence": "parsed" }
```

## WHY THE HASH MATTERS

The excerpt is not decorative. It is re-read from the snapshot and hashed; a fabricated or drifting excerpt fails the locator-support audit.

## WHAT THE RECORD DOES NOT PROVE

- That the excerpt is semantically sufficient.
- That the parser understood every construct.
- That the decision is correct — only that it is inspectable.

## CLAIM-STRENGTH LADDER

### A label caps what may be inferred.

| | | |
|---|---|---|
| **verified-id** | 12 | Public standard ID/title verified; body may be gated. |
| **verified-policy** | 3 | Public policy or hardening guidance; no rule ID claimed. |
| **verified-doc** | 3 | Documented mechanism, e.g. the 24-character APD limit. |
| **secondary** | 3 | Credible secondary source, e.g. an experience report. |
| **heuristic** | 3 | PreFlight pattern; explicitly not a Cloud Standard. |

Counts are the 24 advisory findings in R-DEMO-001. The separate SBT readiness flag is not one of the 24.

## WHAT THE LABEL BUYS YOU

| LABEL | YOU MAY SAY | YOU MAY NOT SAY |
|---|---|---|
| `verified-id` | This is standard ID/title X. | “This passes the gated rule body.” |
| `verified-policy` | Public policy discourages this. | “This is a Cloud Standards violation.” |
| `verified-doc` | The documented mechanism is X. | “The target estate behaves this way.” |
| `secondary` | An experience report flags this. | “This is an official rule.” |
| `heuristic` | PreFlight flags this pattern. | “Guidewire says this is wrong.” |

## ASSUMPTION REGISTER

| | |
|---|---|
| ID | `A-01` |
| Statement | Territory_Ext exists; no rating or territory source artefact was found. |
| Owner | UNASSIGNED |
| Impact | high · heuristic range **3–5 engineer-days** |
| Question | Which rating table or service consumes Territory_Ext, and is it in scope? |
| Status | open |

## PACKET HASH / AUDIT CHAIN

`event 1` → `hash 1` → `event 2` → `hash 2` → … → `head`

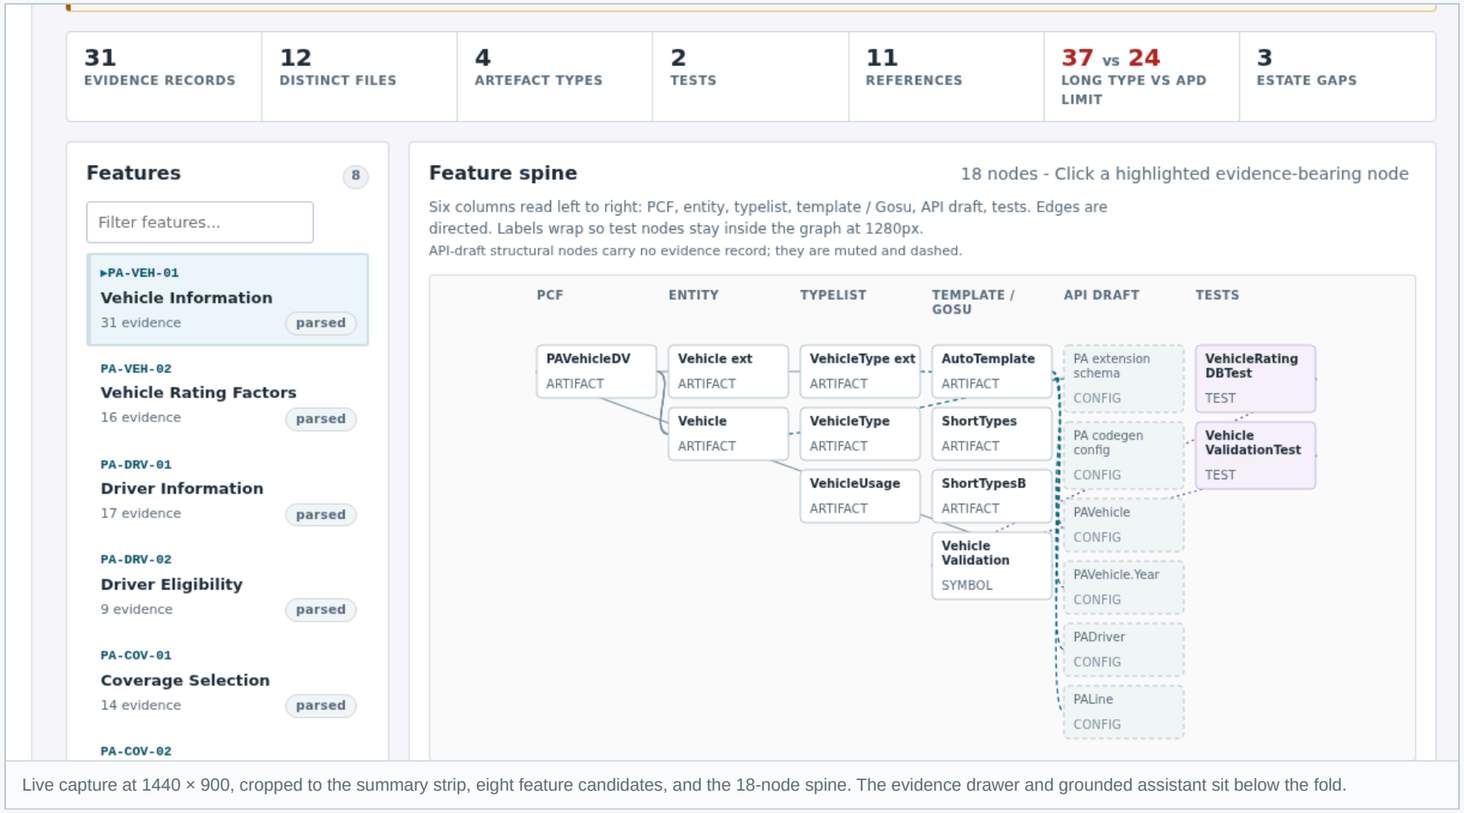
The packet hash covers the rendered packet HTML (hash fields normalized); `packet_json` separately covers canonical packet JSON; content digests bind findings, mappings, and evidence. The audit head covers the append-only event chain. Verification is local and tamper-evident — not non-repudiation.

## REVIEWER QUESTION

If this evidence disappeared, which decision would become unsafe? If the answer is “none,” the record is not carrying its weight.

# One feature, traced across six artefact families.

PA-VEH-01 Vehicle Information is the slice. The summary strip counts what the claim rests on; the spine shows how PCF, entity, typelist, template/Gosu, API draft, and tests connect.




Live capture at 1440 × 900, cropped to the summary strip, eight feature candidates, and the 18-node spine. The evidence drawer and grounded assistant sit below the fold.

1. **31 evidence / 12 files / 4 artefact types / 2 tests / 11 references — all PA-VEH-01 counts.** Every number is a count of linked records. The 37-character long type is separate: it appears in two files (template and product model) and is a PA-VEH-02 dependency, not one of the 11 references.
2. **Eight features, each with evidence count and confidence.** The selected feature is PA-VEH-01; the list is the inventory entry point.
3. **18-node spine, six columns.** PCF → entity → typelist → template/Gosu → API draft → tests. Directed edges; click any highlighted evidence-bearing node for evidence.
4. **37 vs 24, and 3 estate-wide gaps.** The long type is carried by two files and is a PA-VEH-02 dependency; the gaps are labelled estate-wide, not three missing links in this feature.

| 31 | 12 | 37:24 | 3 |
|---|---|---|---|
| EVIDENCE | FILES | TYPE VS APD | ESTATE GAPS |

parsed EV-0001

`config/web/pcf/policy/PAVehicleDV.pcf`

xpath /PCF/Screen/TypeKeyInput[1]:4

```
<TypeKeyInput id="VehicleType" value="VehicleType"
label="VehicleType"/>
```

Method: xml_xpath. Excerpt SHA-256: d76bf9908b1f

**Click node → EV-0001.** Path, xpath, line, excerpt, method, excerpt SHA-256. This is the atomic unit of the review.

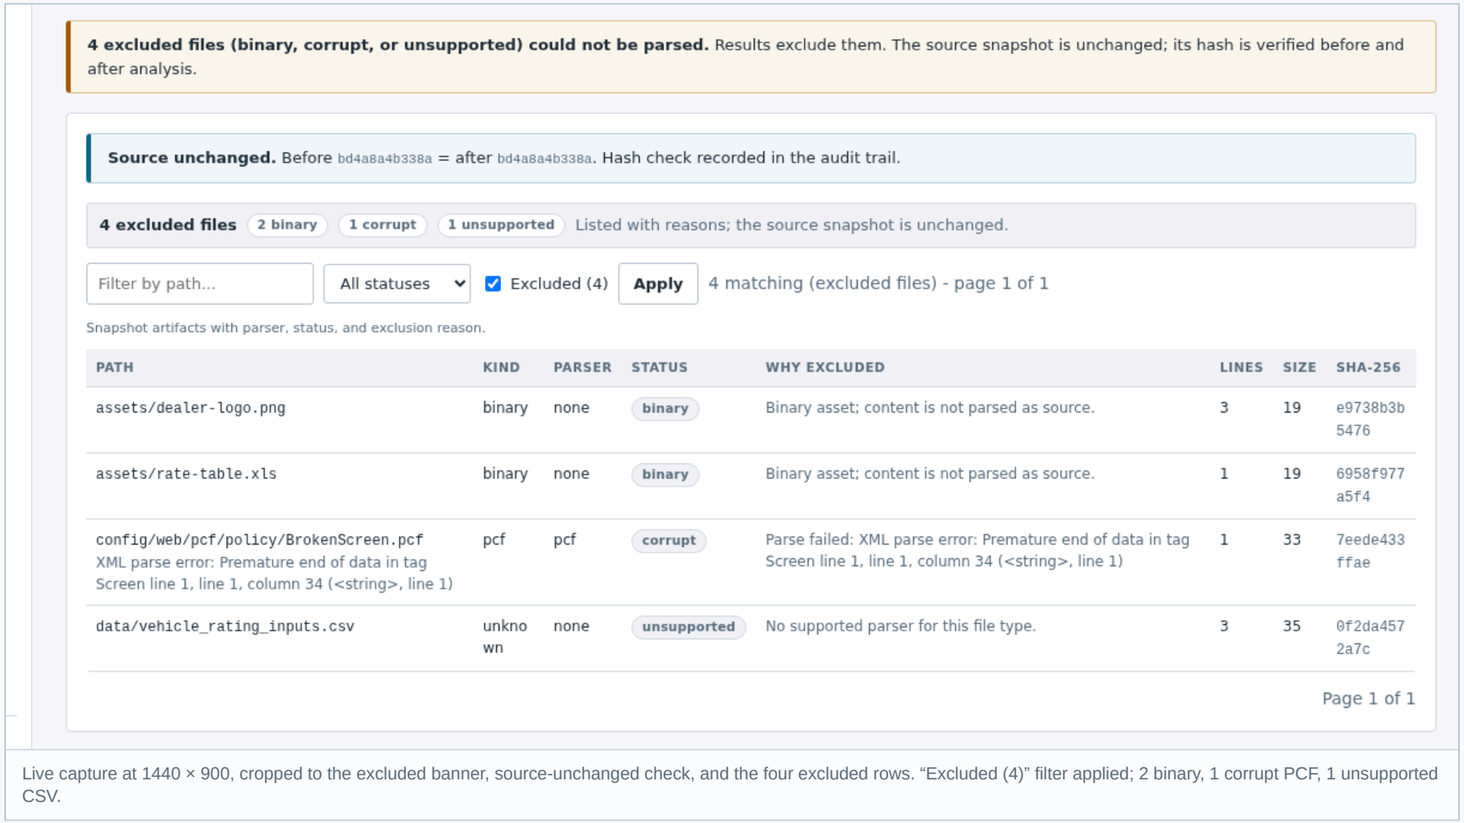
**Fixture realism:** the 300 files include generated filler. PCF, product model, and Gosu are synthetic sketches, not production Guidewire formats. Codegen and entity names match the fixture; this is not a claim of domain fidelity.

DEMO 02 · MANIFEST

300 FILES · 4 EXCLUDED

# No silent omission. Four files are excluded — and named.

A parser that quietly drops what it cannot read produces a confident lie. PreFlight records every file, status, reason, and hash; the source snapshot is checked before and after analysis.



**4 excluded files (binary, corrupt, or unsupported) could not be parsed.** Results exclude them. The source snapshot is unchanged; its hash is verified before and after analysis.

**Source unchanged.** Before `bd4a8a4b338a` = after `bd4a8a4b338a`. Hash check recorded in the audit trail.

**4 excluded files** 2 binary · 1 corrupt · 1 unsupported — Listed with reasons; the source snapshot is unchanged.

Filter by path... | All statuses | Excluded (4) | Apply | 4 matching (excluded files) - page 1 of 1

Snapshot artifacts with parser, status, and exclusion reason.

| PATH | KIND | PARSER | STATUS | WHY EXCLUDED | LINES | SIZE | SHA-256 |
|---|---|---|---|---|---|---|---|
| assets/dealer-logo.png | binary | none | binary | Binary asset; content is not parsed as source. | 3 | 19 | e9738b3b5476 |
| assets/rate-table.xls | binary | none | binary | Binary asset; content is not parsed as source. | 1 | 19 | 6958f977a5f4 |
| config/web/pcf/policy/BrokenScreen.pcf<br>XML parse error: Premature end of data in tag Screen line 1, line 1, column 34 (<string>, line 1) | pcf | pcf | corrupt | Parse failed: XML parse error: Premature end of data in tag Screen line 1, line 1, column 34 (<string>, line 1) | 1 | 33 | 7eede433ffae |
| data/vehicle_rating_inputs.csv | unknown | none | unsupported | No supported parser for this file type. | 3 | 35 | 0f2da4572a7c |

Page 1 of 1

Live capture at 1440 × 900, cropped to the excluded banner, source-unchanged check, and the four excluded rows. “Excluded (4)” filter applied; 2 binary, 1 corrupt PCF, 1 unsupported CSV.

1. **300 files; 4 excluded: 2 binary / 1 corrupt / 1 unsupported.** The manifest is the denominator for every later claim.
2. **Every excluded row has a reason and a SHA-256.** BrokenScreen.pcf fails XML parsing at line 1; `vehicle_rating_inputs.csv` has no supported parser. The filter is part of the UI, not a footnote.
3. **Before = after.** `bd4a8a4b338a` is identical before and after analysis; the immutable snapshot hash is also visible in the live banner. The check is recorded in the audit trail.

| EXCLUDED FILE | STATUS | REASON |
|---|---|---|
| `assets/dealer-logo.png` | binary | not parsed as source |
| `assets/rate-table.xls` | binary | not parsed as source |
| `config/web/pcf/policy/BrokenScreen.pcf` | corrupt | XML parse error, line 1 |
| `data/vehicle_rating_inputs.csv` | unsupported | no supported parser |

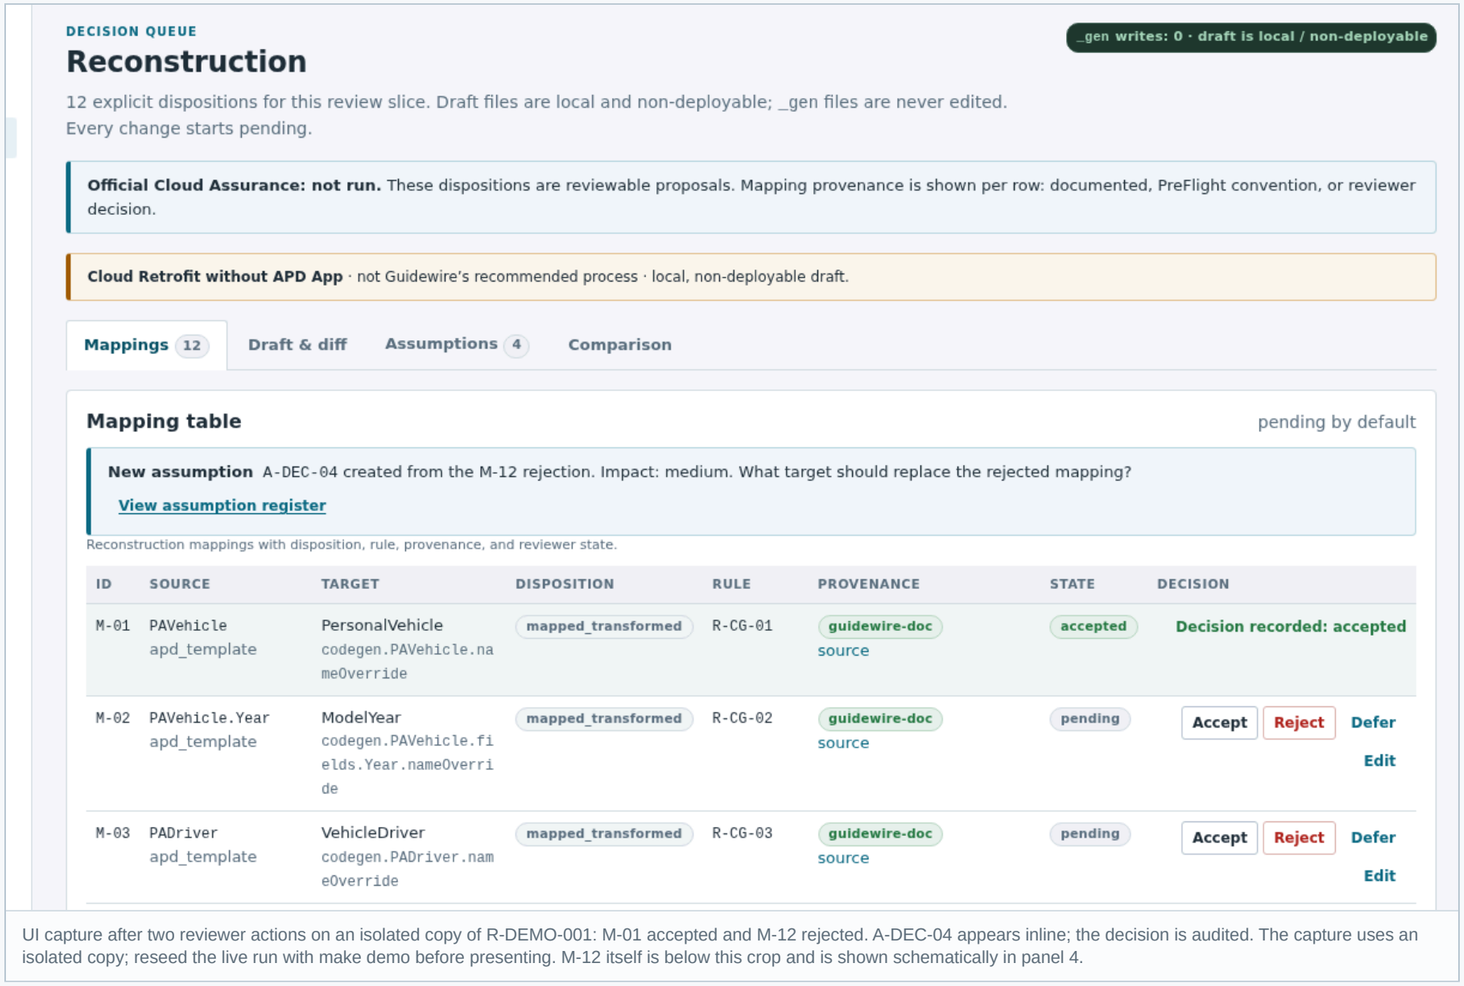
**The discipline:** an inventory without a failure list is a sample. The excluded files stay visible in the packet.

DEMO 03 · RECONSTRUCTION · AFTER REVIEWER ACTIONS

POST-DECISION · 4 ASSUMPTIONS

# A decision queue, not a rename script.

Twelve explicit source-to-target dispositions. Provenance is per row: documented example, PreFlight convention, or reviewer decision. Reject a row and the unresolved question becomes a named assumption.




UI capture after two reviewer actions on an isolated copy of R-DEMO-001: M-01 accepted and M-12 rejected. A-DEC-04 appears inline; the decision is audited. The capture uses an isolated copy; reseed the live run with make demo before presenting. M-12 itself is below this crop and is shown schematically in panel 4.

1. **Reject M-12 → new A-DEC-04.** "What target should replace the rejected mapping?" Impact medium, owner unassigned, status open.
2. **Accept is explicit and recorded in the audit.** M-01 row turns green only after the action; the other rows stay pending.
3. **_gen writes: 0.** Draft is local and non-deployable; only `_ext`, codegen, and product-access proposals are written.
4. **M-12 rejected, not deleted.** 12 dispositions / 5 target changes. Four exact documented examples + one documented 24-char mechanism on a synthetic long type; six API conventions; one typelist convention.

| 12 | 5 | 4 | 0 |
|---|---|---|---|
| DISPOSITIONS | TARGET CHANGES | ASSUMPTIONS AFTER REJECT | _GEN WRITES |

### SCHEMATIC · M-12 ROW BELOW THE CROP

| | |
|---|---|
| ID | M-12 · rejected |
| Source | `VehicleType.Motorcycle_Ext` |
| Target | `schema.vehicleType.motorcycle` |
| State | Stays in the ledger; A-DEC-04 carries the open question. |

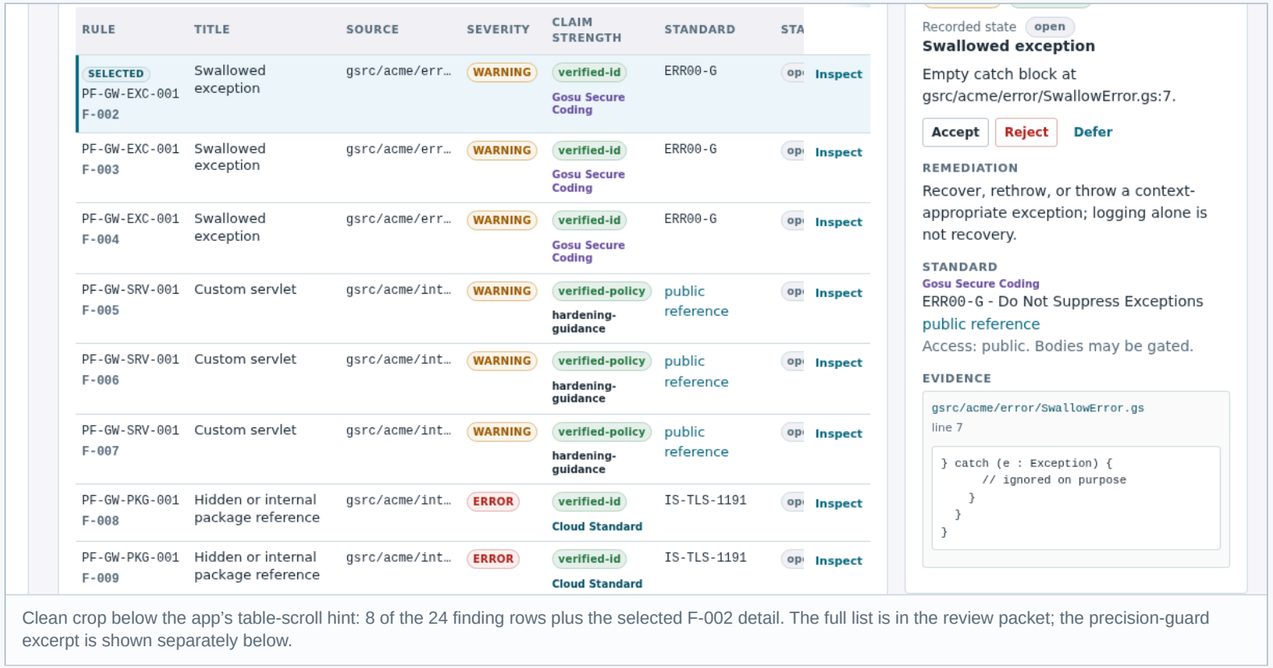
**Scope:** draft targets Cloud Retrofit **without APD App**, which Guidewire explicitly labels not the recommended process; field-level `nameOverride` applies only on that path. The `_ext` draft is a local JSON contract, not official Guidewire schema, mapping, or updater syntax.

DEMO 04 · FINDINGS

OFFICIAL CLOUD ASSURANCE: NOT RUN

# Precision is a feature. Restraint is the interface.

Twenty-four advisory findings from eight rules. Every row carries a claim-strength label. The rule pack ships 24 hard negatives that must stay silent — including the supported `Query.make` API.




Clean crop below the app's table-scroll hint: 8 of the 24 finding rows plus the selected F-002 detail. The full list is in the review packet; the precision-guard excerpt is shown separately below.

1. **24 findings / 8 rules.** 3 per rule; the guard corpus is separate: 8 rules × 3 hard negatives = 24 cases.
2. **Claim strength on every row.** 12 verified-id, 3 verified-policy, 3 verified-doc, 3 secondary, 3 heuristic.
3. **Two standard families, not one.** Cloud Standards: `IS-TLS-1238` (SQL/JDBC), `IS-TLS-1191` (hidden package). Gosu Secure Coding: `SEC00-G` (secrets), `ERR00-G` (swallowed exception). The `standard_family` field in the data distinguishes `cloud-standard` from `gosu-secure-coding`.
4. **SBT is a separate readiness flag.** Not mixed into the 24 advisory findings or the 8 scored rules.
5. **24 hard negatives stay silent.** `Query.make(PersonalVehicle).select()` is supported access, not JDBC. Comments and strings are masked first.

**Precision guard:** 8 rules × 3 near-misses; `Query.make` is supported access, not JDBC. Comments and strings are masked before matching.

**No violation claim.** Advisory, pattern-based, cited. Not Cloud Assurance, a pass/fail, or legal/security advice.

DEMO 05 · GROUNDED ASSISTANT

RETRIEVAL-ONLY · MODEL OFF

# It would rather refuse than invent a source.

Deterministic retrieval over the run's evidence. Model off. A missing rating table produces a refusal with the missing sources named; a present typelist entry produces a citation card with path, locator, and excerpt.

REFUSAL **"Does this estate use territory rating?"**

**You: Does this estate use territory rating?**

> Cannot determine - no rating table or territory source artefact exists in this snapshot. Searched scopes: product model, rating Gosu, rate tables, entity extensions. No evidence in this snapshot supports a different answer.
>
> Would need: a rating table or territory factor source; rating input export; the service that consumes Territory_Ext

1. **Names the missing evidence, not a guess.** "Would need: a rating table or territory factor source; rating input export; the service that consumes Territory_Ext."
2. **Territory is A-01 in the packet.** The refusal and the assumption register point at the same unresolved source.

CITED ANSWER **"Which files reference Motorcycle_Ext?"**

**You: Which files reference Motorcycle_Ext?**

Motorcycle_Ext appears in 1 evidence record, including the typelist extension. The target enum disposition is a reviewer decision, not a source fact. Each citation resolves to a file and line/pointer.

Citations (open the evidence panel)

linked EV-0023 **Open evidence →**

`config/extensions/typelist/VehicleType.ttx`

xpath /typelistExtension/code[1]:3

```
<code name="Motorcycle_Ext"/>
```

3. **A citation card is an evidence record.** `EV-0023` → `config/extensions/typelist/VehicleType.ttx` → xpath `/typelistExtension/code[1]` → line 3.
4. **Disposition is still a reviewer decision.** The assistant reports the source fact; it does not accept the mapping.

**No generation**
No LLM, no tools, no external calls. Retrieval over evidence IDs.

**No orphan citations**
Every cited ID resolves to the embedded evidence record, or is labelled missing.

**No authority transfer**
An answer is evidence to review, not a decision. The human still accepts or rejects.

DEMO 06 · EVALUATION

NOT A BENCHMARK

# Perfect scores measure the instrument. The estate is still unproven.

Six headline metrics on the in-sample demo slice, with raw counts and 95% Wilson intervals. The holdout is pre-computed and never imported. A small n with a 100% point estimate is not certainty.

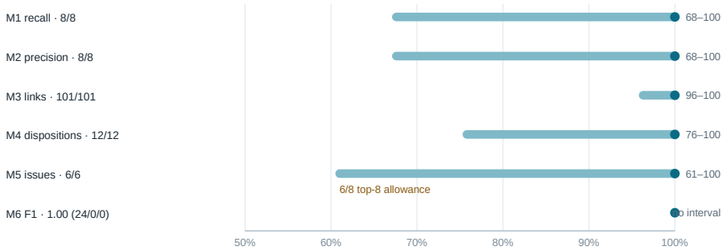
| | | |
|---|---|---|
| **8/8**<br>M1 RECALL | **8/8**<br>M2 PRECISION | **101/101**<br>M3 LINKS + 30/30 SAMPLE |
| **12/12**<br>M4 DISPOSITIONS | **6/6**<br>M5 ISSUES · 6/8 SLOTS | **1.00**<br>M6 F1 · 24 TP / 0 FP / 0 FN |

**Guard count, not a seventh metric:** 24/24 hard negatives silent.

**Holdout:** 4/4 capability recall, 95% Wilson 51.0%–100.0%, four-feature naming-shift fixture. **Not imported** into the demo or used for tuning.

## 95% Wilson intervals

dot = observed value · line = interval · M6 has no interval (bootstrap deferred)




Wilson score intervals from reported numerators and denominators. M6 has no interval at this n; bootstrap deferred. Demo metrics are in-sample on one synthetic fixture. 24/24 hard negatives are a separate guard count.

### PRE-REGISTERED MARGINS VS BASELINES

| METRIC | BASELINE | DELTA | REQUIRED | STATUS |
|---|---|---|---|---|
| M1 feature recall | PCF filename baseline: 87.5% | +12.5 pp | +10 pp | PASS |
| M4 dispositions | Identity mapper: 50.0% | +50.0 pp | +10 pp | PASS |
| M6 finding micro-F1 | Grep-only F1: 84.0% | +16.0 pp | +10 pp | PASS |

### WHAT THE NUMBERS DO AND DO NOT SAY

- (+) **M3 is 101/101 all links** plus a 30/30 stratified sample. It checks excerpt support at the source, not semantic entailment.
- (+) **M5 is 6/6** but the allowance is **6 of 8 slots**. It is a fixed two-decoy top-8, not a learned ranking.
- (+) **M6 is F1 1.00 (24 TP / 0 FP / 0 FN)** = 48/48 micro-F1. Line-level matching is deferred; the demo uses path-level matching.
- (–) **No human IAA.** The reference target is internally authored and AI-assisted, not blind and independent.
- (–) **One synthetic fixture.** The four-feature holdout is naming-shift only and never imported or tuned.
- (–) **Not a benchmark.** No real estate, customer code, deployment, or official Cloud Assurance run.

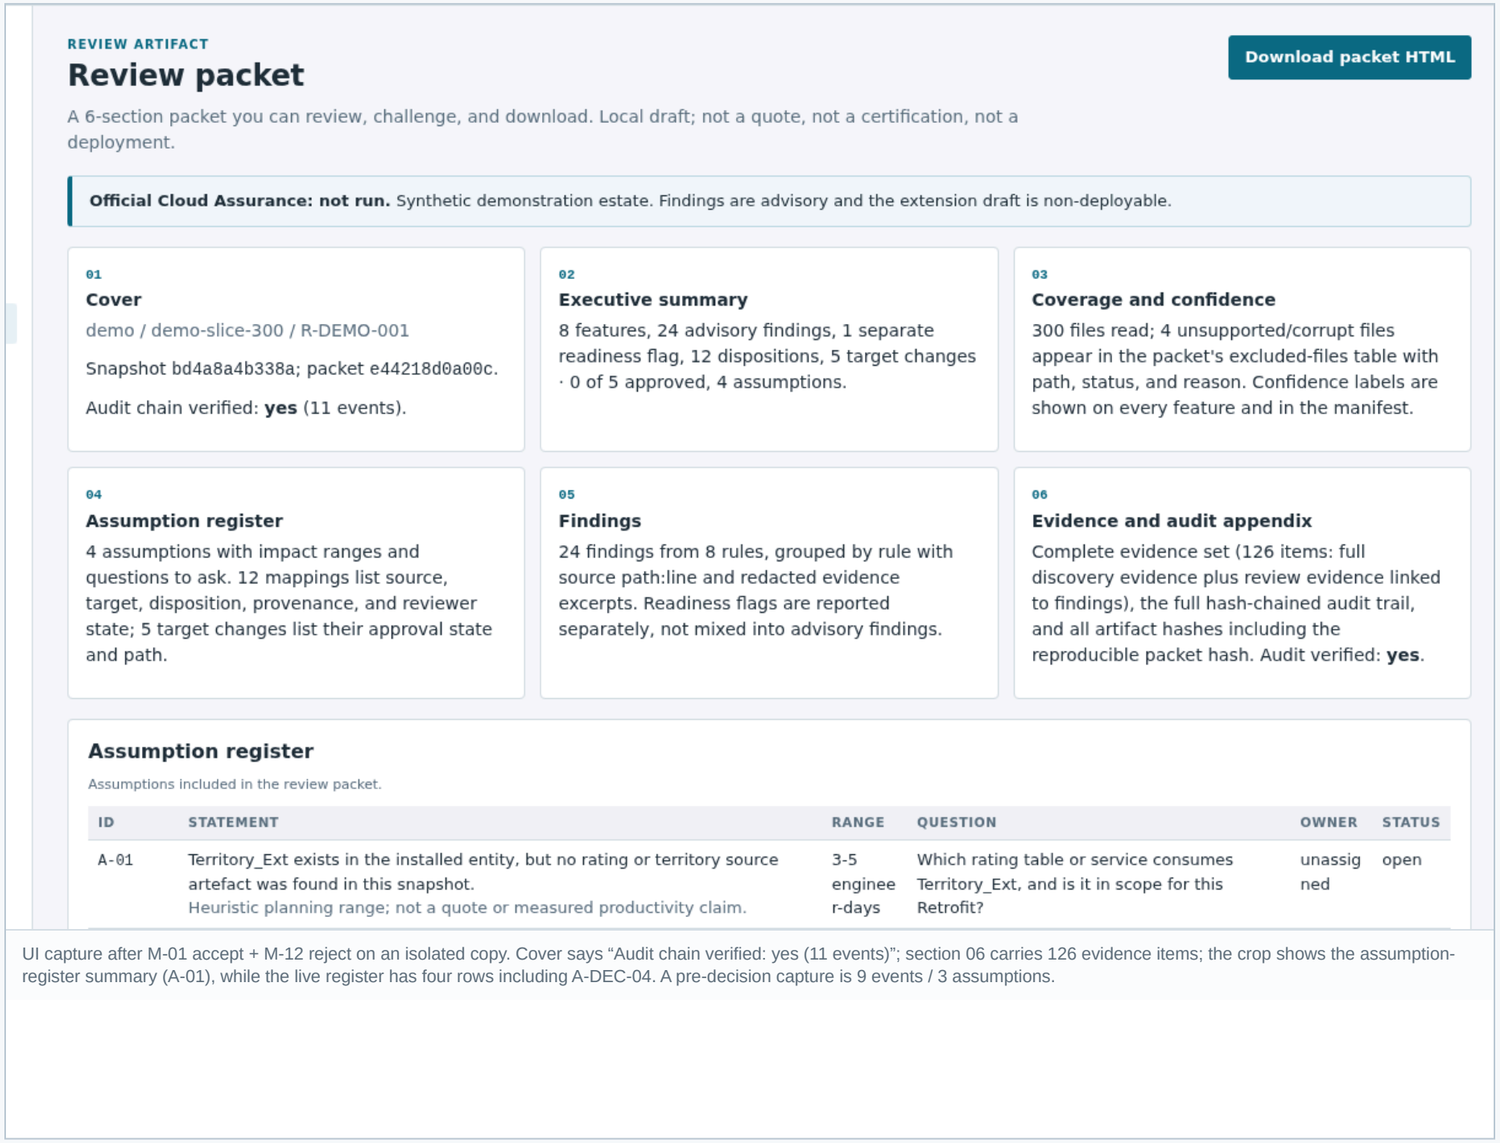
### WHAT WOULD CHANGE THE VERDICT

- A real anonymised estate and a second LOB.
- Blind human annotation with reported IAA.
- Line-level finding evaluation and compile validation.

POST-DECISION · 11 EVENTS · 4 ASSUMPTIONS

# The review artifact is the product.

Six sections, every excluded file, every assumption, every disposition, 126 evidence items, and a hash-chained audit trail. This capture is post-decision: 11 audit events and 4 assumptions after Accept M-01 + Reject M-12. Downloadable HTML; locally verifiable; explicitly not a certification.



REVIEW ARTIFACT

## Review packet

Download packet HTML

A 6-section packet you can review, challenge, and download. Local draft; not a quote, not a certification, not a deployment.

**Official Cloud Assurance: not run.** Synthetic demonstration estate. Findings are advisory and the extension draft is non-deployable.

01
**Cover**
demo / demo-slice-300 / R-DEMO-001

Snapshot bd4a8a4b338a; packet e44218d0a00c.

Audit chain verified: **yes** (11 events).

02
**Executive summary**
8 features, 24 advisory findings, 1 separate readiness flag, 12 dispositions, 5 target changes · 0 of 5 approved, 4 assumptions.

03
**Coverage and confidence**
300 files read; 4 unsupported/corrupt files appear in the packet's excluded-files table with path, status, and reason. Confidence labels are shown on every feature and in the manifest.

04
**Assumption register**
4 assumptions with impact ranges and questions to ask. 12 mappings list source, target, disposition, provenance, and reviewer state; 5 target changes list their approval state and path.

05
**Findings**
24 findings from 8 rules, grouped by rule with source path:line and redacted evidence excerpts. Readiness flags are reported separately, not mixed into advisory findings.

06
**Evidence and audit appendix**
Complete evidence set (126 items: full discovery evidence plus review evidence linked to findings), the full hash-chained audit trail, and all artifact hashes including the reproducible packet hash. Audit verified: **yes**.

### Assumption register

Assumptions included in the review packet.

| ID | STATEMENT | RANGE | QUESTION | OWNER | STATUS |
|---|---|---|---|---|---|
| A-01 | Territory_Ext exists in the installed entity, but no rating or territory source artefact was found in this snapshot. Heuristic planning range; not a quote or measured productivity claim. | 3-5 engineer-days | Which rating table or service consumes Territory_Ext, and is it in scope for this Retrofit? | unassigned | open |

UI capture after M-01 accept + M-12 reject on an isolated copy. Cover says “Audit chain verified: yes (11 events)”; section 06 carries 126 evidence items; the crop shows the assumption-register summary (A-01), while the live register has four rows including A-DEC-04. A pre-decision capture is 9 events / 3 assumptions.

1. **Downloadable HTML.** A local, self-contained review artifact — not a dashboard link that expires.
2. **Six sections:** cover · executive summary · coverage and confidence · assumption register · findings (24 + 1 separate readiness flag) · evidence and audit appendix.
3. **126 evidence items** = full discovery evidence plus review evidence linked to findings. Not a 20-row sample.
4. **Audit verified: yes, 11 events post-decision.** Seed state is 9 events; Accept M-01 + Reject M-12 add two. The packet hash seals the rendered packet HTML; `packet_json` seals canonical JSON; content digests bind findings, mappings, and evidence; the audit head seals the event chain. Two-step verification: `verify-packet` checks the packet and content digests; `verify-audit` checks the audit chain only, not the packet. Local tamper-evident, not non-repudiation.

| | |
|---|---|
| 01 Cover + hashes | 02 Executive summary |
| 03 Coverage & confidence | 04 Assumption register |
| 05 Findings | 06 Evidence & audit |

12 mappings · 4 assumptions · 1 readiness flag · 4 excluded files · 11 audit events

```
packet      read live rendered packet HTML
packet_json read live canonical JSON
audit head  read live event-chain head
snapshot    bd4a8a4b338a… fixture; stable
```

```
$ preflight verify-packet R-DEMO-001
```

```
$ preflight verify-audit R-DEMO-001
```

Content digests bind findings, mappings, and evidence. Packet, packet_json, and audit-head values change with reviewer actions and packet refreshes — read them from the live screen. `verify-packet` checks the packet; `verify-audit` checks the audit chain only, not the packet. The bundled packet is a capture-time example, not a live artifact.

DEMO 08 · MIGRATOR · FIVE-AGENT HANDOFF CHAIN

SEEDED RUN · 3 ASSUMPTIONS

# Five agents. One handoff chain. No authority to deploy.

The feature migrator runs five specialist agents over the completed PA-VEH-01 review: Scout, Architect, Guardian, Critic, and Scribe. Each receives the prior artifact, adds a check or an artifact, and leaves a record. The run ends `completed_with_gates` with 11 gates open — not approval, not deployment. This path is a straight handoff chain; the codebase lab is the cyclic loop.

**Feature handoff chain:** Scout → Architect → Guardian → Critic → **Scribe** seals the review packet. **Codebase cyclic loop:** Scout → Architect ↔ Guardian / Critic → **Planner** assembles and seals the codebase plan.

**Same fifth role, different loop:** seal the run, never approve or deploy.

01
### Scout
**REVIEW → SCOPE**

300 files · 101 evidence records · 31 links on PA-VEH-01 · 24 findings · 12 dispositions

→

02
### Architect
**SCOPE → DRAFT**

5 work packages · 5 target files · all `draft_ready` · all `pending`

→

03
### Guardian
**DRAFT → GATES**

11 open gates = 3 assumptions + 8 rule groups · 6 high impact

→

04
### Critic
**GATES → CHECKS**

7/7 checks passed · verdict PASS · 0 auto-approved changes

→

05
### Scribe
**CHECKS → SEAL**

8 listed bundle files · 5 targets + plan + report + 12-event log · manifest seals them · `completed_with_gates`

**BOUNDED AUTONOMY**

Read the snapshot, generate local drafts, verify, seal. Never deploy, edit _gen, approve a change, resolve an assumption, or contact a tenant.

**REALITY BOUNDARY**

**Synthetic rehearsal, not an actual migration.** No tenant contacted. No Gosu/Java compiled against Guidewire. No official Cloud Assurance result.

**R-DEMO-001** · **completed_with_gates** · verification passed (7/7) · 11 gates open · no deployment

**State:** the seeded review has 3 assumption gates + 8 rule groups = 11 gates. Slides 8 and 12 show the separate post-decision packet (4 assumptions after Reject M-12); that is not the migrator's input state. Optional live reasoning run `R-20261111-052740-b001`: `mimo-v2.6-flash` made 3 citation-checked calls (Scout, Architect, Critic); 0 ungrounded claims passed validation; Guardian made 0 model calls. The deterministic 7/7 checks and sealed status are unchanged.

## PLAIN-LANGUAGE GLOSSARY

**GOSU**
Guidewire's JVM language for business logic and queries.

**APD**
Advanced Product Designer: product data and code generation.

**SCRIBE**
The feature handoff chain's fifth agent; seals the review packet.

**PLANNER**
The codebase cyclic loop's fifth agent; assembles and seals the plan.

5 PACKAGES · 12 DISPOSITIONS

# Scout hands over 31 links and 12 dispositions. Architect writes five files.

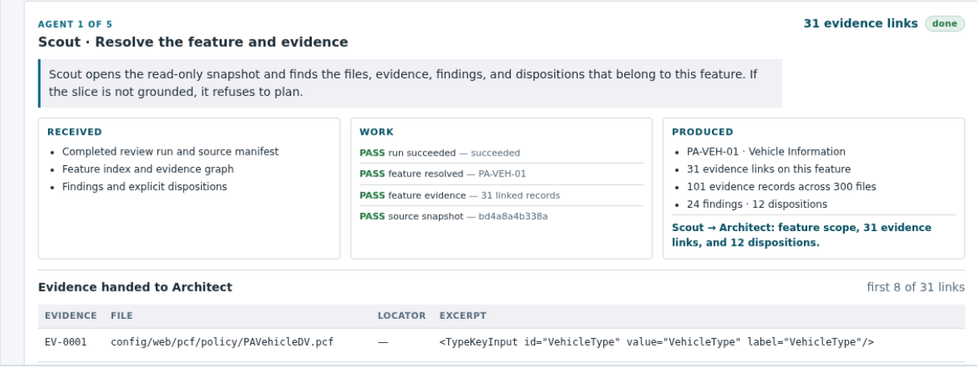
Scout resolves PA-VEH-01 from 300 files and 101 evidence records, keeping 31 links, 24 findings, and 12 dispositions attached to the handoff. Architect turns those dispositions into five local targets: schema, mapping, updater, codegen config, and API contract. The mapping file carries all 12 dispositions, so five files do not mean five trivial changes. Every target is `draft_ready` and `pending`.

## Scout → Architect

300 files · 101 evidence records · 31 links on PA-VEH-01 · 24 findings · 12 dispositions · snapshot `bd4a8a4b338a`

| WP | TARGET (BASENAME) | DISPOSITION | EVIDENCE | BYTES | GATE / APPROVAL |
|---|---|---|---|---|---|
| WP-01 | policyperiod_pa_ext-1.0.schema.json | PAVehicle → PersonalVehicle | EV-0011, EV-0018 | 945 B | — / pending |
| WP-02 | policyperiod_pa_ext-1.0.mapping.json | PAVehicle.Year → ModelYear | EV-0002, EV-0006 | 3,046 B | — / pending |
| WP-03 | policyperiod_pa_ext-1.0.updater.json | PADriver → VehicleDriver | EV-0057, EV-0061 | 402 B | — / pending |
| WP-04 | pa_codegen_config_ext-1.0.yaml | PALine → PersonalAutoLine | EV-0097 | 520 B | — / pending |
| WP-05 | pa_ext-1.0.yaml | PersonalAutoRatingFactorConfiguration → same | EV-0036 | 78 B | **G-03** / pending |

Full target paths are under `migration/target/config/integration/`. All five are `mapped_transformed`, `draft_ready`, and `pending`.




Scout stage on `R-DEMO-001`: the crop shows the first evidence row of 31 on PA-VEH-01; the full handoff is in the packet. The package table above is built from `migration.json`, not from this screenshot.

### FIVE FILES, TWELVE DISPOSITIONS

WP-02 is a 12-entry mapping file; WP-05 is the only package carrying a gate at handoff (G-03, 37-character APD type). Five target files are not five deployment units.

**_gen** targets: 0 · duplicate targets: 0 · auto-approved: 0 · all five `pending`

Work package = one target file generated from one explicit disposition. APD = Advanced Product Designer; the 37-character type exceeds the 24-character type-name import limit and needs a codegen `nameOverride`.

DEMO 10 · MIGRATOR · GUARDIAN

SEEDED RUN · 3 ASSUMPTIONS · 11 GATES

# Guardian turns 24 findings and 3 assumptions into 11 human gates.

Guardian does not score the draft; it opens a gate for every open assumption and every rule group that still has an unanswered finding. The 24 findings collapse into 8 rule groups of 3, plus 3 assumption gates, for 11 human decisions. Six of the 11 are high impact. The bundle cannot ship while any gate is open.

**State:** seeded review — 3 assumptions + 8 rule groups = 11 gates. Slides 8 and 12 show the separate post-decision packet (4 assumptions after Reject M-12).

| 24 | 8 | 3 | 11 | 6 | 0 |
|---|---|---|---|---|---|
| FINDINGS | RULE GROUPS × 3 | ASSUMPTIONS | HUMAN GATES | HIGH IMPACT | RESOLVED |

AGENT 3 OF 5

**Guardian · Open deterministic risk gates**

11 gates done

Guardian is the strict compliance gate. It sees only the Architect's packages and evidence-linked findings/assumptions. It opens a gate for every open assumption and every rule family that still has an unanswered finding.

**RECEIVED**

- Architect's 5 packages
- Open assumptions and review findings
- Guidewire rule pack and APD constraints

**DELIBERATELY NOT RECEIVED**

- Scout's narrative or opinion
- Critic's verdict or suggested fixes
- Any model-generated text

**WORK**

PASS assumption gates — 3 open

PASS finding gates — 8 rule groups

PASS no self-approval — all changes remain pending

**PRODUCED**

- 11 open gates: 3 assumption + 8 rule groups
- 6 high-impact gates
- Deployment remains blocked

**Guardian → Critic: 11 gates; deterministic verdict is independent of the narratives.**

**Open gates by type**

3 assumptions · 8 rule groups

**ASSUMPTION GATES (3)**

G-01 high

**Territory_Ext exists in the installed entity, but no rating or territory source artefact was found in this snapshot.**

Which rating table or service consumes Territory_Ext, and is it in scope for this Retrofit?

Owner: unassigned

G-02 medium

**Rating inputs were not in a supported export, so rating behaviour cannot be traced from this snapshot.**

Can the team provide a supported rating input export before scope is agreed?

Owner: unassigned

**RULE GROUPS (8)**

| RULE | FINDINGS | SEVERITY | GATE |
|---|---|---|---|
| PF-GW-EXC-001 | 3 | medium | G-04 |
| PF-GW-GUNIT-001 | 3 | medium | G-05 |
| PF-GW-NAM-001 | 3 | high | G-06 |
| PF-GW-PKG-001 | 3 | high | G-07 |
| PF-GW-SEC-001 | 3 | high | G-08 |
| PF-GW-SQL-001 | 3 | high | G-09 |
| PF-GW-SRV-001 | 3 | medium | G-10 |

## 3 ASSUMPTION GATES

| GATE | DECISION NEEDED | OWNER |
|---|---|---|
| G-01 HIGH | Which rating table or service consumes `Territory_Ext`, and is it in Retrofit scope? | unassigned |
| G-02 MED | Can the team provide a supported rating input export before scope is agreed? | unassigned |
| G-03 HIGH | Who owns the 37-character `PersonalAutoRatingFactorConfiguration` change, and can the `nameOverride` be validated? | unassigned |

## 8 RULE GROUPS · 24 FINDINGS

| RULE GROUP | FINDINGS | IMPACT | GATE |
|---|---|---|---|
| PF-GW-EXC-001 | 3 | medium | G-04 |
| PF-GW-GUNIT-001 | 3 | medium | G-05 |
| PF-GW-NAM-001 | 3 | high | G-06 |
| PF-GW-PKG-001 | 3 | high | G-07 |
| PF-GW-SEC-001 | 3 | high | G-08 |
| PF-GW-SQL-001 | 3 | high | G-09 |
| PF-GW-SRV-001 | 3 | medium | G-10 |
| PF-GW-TX-001 | 3 | medium | G-11 |

**Level check:** the 11 gates live at assumption and rule-group level, not package level. Slide 14 shows the only package-level gate, G-03.

**Independent by design**

Guardian receives the 5 packages, open assumptions and findings, and the rule pack. It deliberately does not receive Scout's narrative, Critic's verdict, or any model-generated text.

**Open is a status, not a failure**

A gate is a named human decision with an owner. All 11 are unresolved; 6 are high impact.

DEMO 11 · MIGRATOR · CRITIC → SCRIBE

7/7 CHECKS · 8 FILES SEALED

# Critic passes 7/7. Scribe seals 8 listed files. Neither approves.

Critic re-checks the handoff independently: packages exist, paths are safe and unique, generated JSON/YAML parses, every package is evidence-linked, and nothing was auto-approved. Scribe writes the plan, report, manifest, and 12-event log; the manifest hashes the bundle and the log records the handoffs. Sealed means tamper-evident, not approved. The status remains `completed_with_gates`, with all five changes pending.

## CRITIC · 7 MACHINE CHECKS

| # | CHECK | RESULT |
|---|---|---|
| 1 | packages exist | **PASS** · 5 work packages |
| 2 | all package files generated | **PASS** · 5/5 |
| 3 | no _gen targets | **PASS** · bundle is _ext / codegen / updater only |
| 4 | no duplicate targets | **PASS** · 5 unique |
| 5 | generated JSON/YAML parses | **PASS** · 5 files parsed |
| 6 | every package is evidence-linked | **PASS** · 5/5 linked |
| 7 | human approval retained | **PASS** · 0 auto-approved changes |

## SCRIBE · FOUR SEAL CARDS

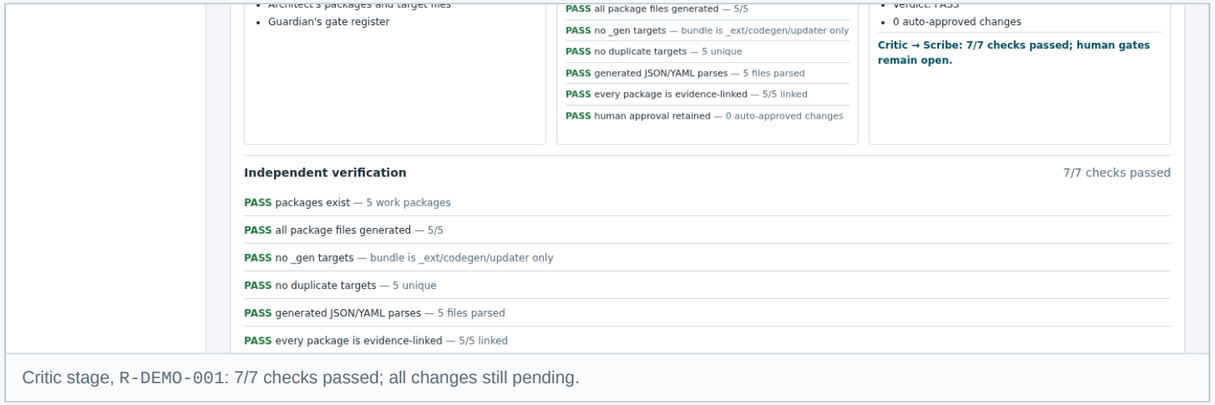
**PLAN**
`3c307f75205c`
`migration-plan.json`

**REPORT**
`8770963bd2fd`
`MIGRATION-REPORT.md`

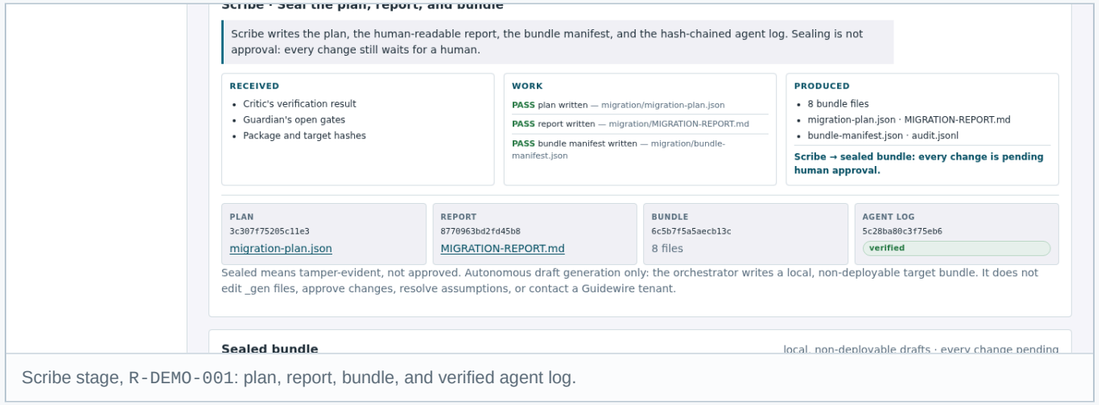
**BUNDLE**
`6c5b7f5a5aec`
8 listed files

**AGENT LOG**
`5c28ba80c3f75eb6`
12 events · **verified**

8 listed bundle files = 5 target drafts + plan + report + 12-event log. `bundle-manifest.json` is the seal wrapper. Sealed ≠ approved.




Critic stage, `R-DEMO-001`: 7/7 checks passed; all changes still pending.




Scribe stage, `R-DEMO-001`: plan, report, bundle, and verified agent log.

Review packet audit: **9 seed events** (**11** after the two runbook decisions). Migrator log: **12 events**, a separate hash chain. Sealed is tamper-evident, not non-repudiation.

**R-DEMO-001** · **completed_with_gates** · verification passed (7/7) · 11 gates open · 0 auto-approved · no deployment

DEMO 12 · CODEBASE LAB · TWO MIGRATION SHAPES

BOTH GREEN_WITH_HUMAN_GATES

# Two synthetic insurers, two migration shapes, one bounded crew.

Cascade Mutual is a monoline Personal Auto mutual; Harbor P&C is a multi-line carrier spanning Personal Auto, Homeowners, and Workers' Comp. Both run the same five-agent codebase loop against synthetic Qusar 2026.08 snapshots with a simulated GWCP 2027.04 target — not a published release. The numbers differ because the estates differ, not because the method changes. Both end `green_with_human_gates`, with every patch pending.

Qusar 2026.08 pre-cloud snapshot → **simulated** GWCP 2027.04 target · five-agent loop Scout → Architect ↔ Guardian / Critic → Planner · same 16-family rule pack · SARIF 2.1.0 · every patch `pending` · no deployment

Same fifth role, different loop: **Scribe** seals the feature packet; **Planner** assembles and seals the codebase plan. Neither approves.

MONOLINE PERSONAL AUTO · SYNTHETIC

RETROFIT WITHOUT APD APP

## Cascade Mutual

**Illustrative synthetic footprint:** 420,000 policies · 38 PCF screens · 64 custom Gosu files · 7 integrations.

Narrow surface: telematics rating, dealer feed, commission publication. Direct SQL in the rating path.

| 18 | 22 | 2 | 10 |
|---|---|---|---|
| FINDINGS<br>10 ERROR / 8 WARNING | PATCHES | ITERATIONS | HUMAN GATES |

`green_with_human_gates` · 13 rules fired · structural 100 / gate-adjusted **74**.

MULTI-LINE REGIONAL P&C · SYNTHETIC

TYPE-LEVEL APD NAMING PREPARATION

## Harbor P&C

**Illustrative synthetic footprint:** 2.6M policies · 214 PCF screens · 1,180 custom Gosu files · 42 integrations.

Cross-application: coastal eligibility, claims gateway, vendor MQ, billing batch.

| 20 | 21 | 2 | 11 |
|---|---|---|---|
| FINDINGS<br>11 ERROR / 9 WARNING | PATCHES | ITERATIONS | HUMAN GATES |

`green_with_human_gates` · 15 rules fired · structural 100 / gate-adjusted **71**.

**Not full APD conversion:** type-level naming preparation only; no field mapping or endpoint regeneration.




Outcome strips only, cropped above the app's program label: Cascade 18/22/2/10 and Harbor 20/21/2/11. The app says "blockers found" and "patches drafted"; the cards above call the same counts findings and patches. Footprints are illustrative synthetic values.

**Reality boundary:** synthetic profiles, not customer references. No customer data, no tenant connection, no Guidewire affiliation. A code-level green is bounded; human gates remain open. Harbor's APD draft is type-level APD naming preparation only — not full APD conversion or field mapping.

# The code lab closes its loop. Cascade 10 and Harbor 11 human gates stay open.

The codebase lab starts from raw Gosu, Java, XML, and GUnit rather than a completed review. Architect drafts patches, Guardian re-scans the target, Critic applies every displayed diff, and Planner assembles and seals the codebase plan. Nothing is compiled: compile validation is future work. Cascade and Harbor each need two passes, and every human gate remains open. The output is a review bundle and SARIF, not a deployed change.

**16** code rule families · 2 cases · 2 iterations each · **Cascade** 18 findings / 22 patches / 10 gates · **Harbor** 20 findings / 21 patches / 11 gates · 13 / 15 SARIF rules fired

### BEFORE · FIXTURE SOURCE (EXCERPT)

```
uses java.sql.Connection
uses java.sql.DriverManager

var conn = DriverManager.getConnection(
  "jdbc:cascade:rating")
var stmt = conn.createStatement()
stmt.executeQuery(
  "SELECT factor FROM cascade_rating " +
  "WHERE vin = '" + vin + "'")
```

Source excerpt from `gsrc/cascade/rating/VehicleRatingQuery.gs`; line breaks added for display. This is the fixture's JDBC and string-concatenated SQL.

### AFTER · CORRECTED DRAFT SHAPE (NOT COMPILED)

```
package cascade.rating

uses gw.api.database.Query

// Architect draft: typed Gosu Query / API service.
// Confirm entity and fields with the data owner.
class VehicleRatingQueryService {
  function load() : List<VehicleRating> {
    return Query.make(VehicleRating).select().toList()
  }
}
```

The app generator currently appends `QueryService` to the source class name, producing `VehicleRatingQueryQueryService`; that naming artifact is owned by the app team, so the corrected name is shown here. The draft is not compiled; semantic equivalence is gate CR-03.

### SARIF 2.1.0 · ADVISORY BY CONSTRUCTION

**SARIF is the machine-readable output of a static analyzer** — findings with rule IDs and file locations. It is advisory input to CI, not a compliance verdict.

Tool driver **PreFlight CodeGuard**. Cascade: **18** results / **13** rules. Harbor: **20** results / **15** rules.

Per result: `ruleId` · `level` · artifact URI · start line/column · `partialFingerprints.primaryLocationLineHash` · claim strength · remediation · redacted flag.

Run properties: `official_cloud_assurance: false` · `synthetic: true`. No source excerpts or secret values; works with `upload-sarif`.

### WHAT THE DRAFT DOES NOT PROVE

- Compile or runtime equivalence.
- Entity or field-level mapping.
- Tenant compatibility or deployment.

**16 rule families.** 13 fired on Cascade, 15 on Harbor; 18 vs 20 SARIF results; both cases 2 passes. **Planner assembles and seals the codebase plan; nothing is compiled.** The lab is a structural scanner, not a compiler or data-flow engine. Compile validation is future work; every patch remains `pending`.

### PLAIN-LANGUAGE GLOSSARY

**GOSU**
Guidewire's JVM language for business logic and queries.

**APD**
Advanced Product Designer: product data and code generation.

**CDA**
Cloud Data Access: streams/exports to S3 Parquet; not a synchronous read.

**APP EVENTS**
Platform message envelopes; the simulated v2 adds tenant and event IDs.

**SARIF**
Machine-readable static-analysis results; advisory, not a compliance verdict.

DEMO 14 · CODEBASE LAB · UPGRADE RESILIENCE

5 SIMULATED SCENARIOS

# A green signal is release-bound, not permanent.

A migration assessment can outlive the release it was written against. PreFlight simulates five GWCP release changes and invalidates the prior green signal rather than silently reusing it. It identifies affected patches, runs scenario checks, and records one of three resume choices: revalidate, rollback, or defer. The APD scenario reruns at its simulated 20-character limit; the other four remain `revalidation_required` until their drafts are fixed.

On detection: assessment generation **1 → 2** · previous green invalidated · affected patches identified · every patch still `pending` · **simulated release scenario**, not a published 2027.04 release

| SIMULATED RELEASE CHANGE | EFFECTIVE CHANGE | AFFECTED RULES | CASCADE PATCHES | HARBOR PATCHES |
|---|---|---|---|---|
| **APD naming** (simulated) | 24 → 20 characters | `PF-CODE-APD-001` | 2 | 2 |
| App Events v2 | envelope adds `tenantId` / `eventId` | `PF-CODE-MSG-001` | 3 | 3 |
| Gateway mTLS | mTLS + scoped service identity | `PF-CODE-SRV-001, PF-CODE-MSG-001` | 7 | 6 |
| Removed Gosu API | deprecated runtime API family removed | `PF-CODE-API-001, PF-CODE-API-002` | 14 | 14 |
| Cloud API version | v1 → v2 route prefix | `PF-CODE-SRV-001` | 4 | 3 |

**Per-scenario exposure — not additive.** Each row is an independent release change; one patch can appear in more than one row. Cascade has 30 patch-exposures across 22 patches (2+3+7+14+4); Harbor has 28 across 21 (2+3+6+14+3). The estate totals stay 22 and 21.

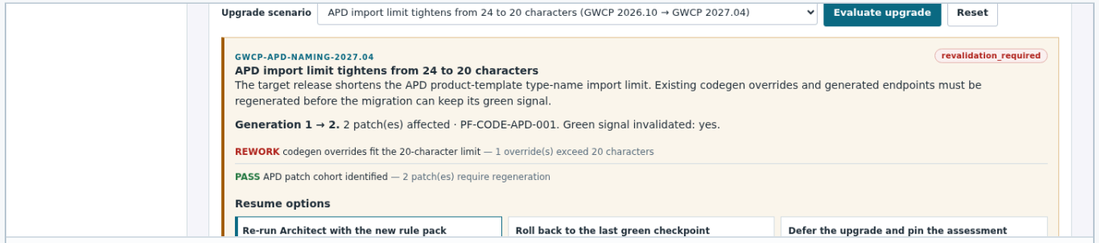
Scenario IDs: `GWCP-APD-NAMING-2027.04` · `GWCP-APP-EVENTS-V2-2027.04` · `GWCP-GATEWAY-MTLS-2027.04` · `GWCP-GOSU-API-REMOVAL-2027.04` · `GWCP-CLOUD-API-VERSION-2027.04`. All five simulate GWCP 2026.10 → 2027.04; the 2027.04 release and the 20-character limit are PreFlight simulations, not published Guidewire release notes.

## RESUME OPTIONS · REVALIDATE / ROLLBACK / DEFER

| OPTION | WHAT IT DOES | STATUS RECORDED |
|---|---|---|
| **Revalidate** RECOMMENDED | Re-run Architect with the new rule pack | `resume_requested` |
| **Rollback** | Return to the last green checkpoint; keep the sealed bundle | `rolled_back_to_checkpoint` |
| **Defer** | Pin the assessment release; record the risk owner | `upgrade_deferred` |

All three clear the green signal and keep human approval required.

**APD control:** same bundle, two outcomes — regenerated at the scenario's 20-character limit → `compatible`; the 24-character bundle → `revalidation_required`.




Upgrade tab: generation bump, invalidated green signal, and the three recorded resume choices.

**Unexpected-scenario playbook**
Named states, not an unchecked green signal: source changed · corrupt file · duplicate target · analyzer crash · empty snapshot.

**Reality boundary**
Conservative match by affected rule or path. Counts are exposure, not effort or defect counts. A human still approves every patch.

ENGINEERING RIGOR

INTERNAL ADVERSARIAL REVIEW

# The demo is the surface. The hardening is the work.

**481 tests** and a Playwright golden path. **18/18** recorded freeze hashes recompute. A security contract with archive, XML, YAML, CSRF, CSP, and redaction tests. Adversarial review found real flaws; each row below was fixed and regression-tested.

**481** tests passed | **1** Playwright golden path | **18/18** frozen hashes recompute | **0** raw secret values in artifacts | **9** seed audit events verified

| FLAW FOUND | BEFORE | AFTER / CONTROL |
|---|---|---|
| **Concurrent mapping decisions** | Simultaneous accepts could lose writes | 12/12 persisted; per-run lock + transaction; 409 on conflict |
| **Secret leakage through evidence** | Parser warnings / adjacent context could store raw values | Redact at capture; fingerprint only; packet and logs tested |
| **YAML alias bomb** | Recursive aliases could expand without bound | Aliases rejected before expansion; depth/node caps; sub-second test |
| **Forged evidence scoring** | Fabricated excerpts could still match gold paths | 48/48 → 0/48 when excerpts are forged; locator support required |
| **Packet evidence completeness** | A 20-row discovery sample | 126 complete evidence items + audit appendix |
| **Evidence audit scope** | A 30-link stratified sample was the headline | 101/101 all links + 30/30 sample; fabricated excerpts fail |
| **Audit verification cost** | Repeated chain scans and no truncation checkpoint | Single-pass walk + checkpoint head; 9 seed events verified |
| **Mobile layout** | Horizontal overflow at 390 px | Clipped overflow + responsive regression coverage |

127.0.0.1 BIND | HOST / ORIGIN ALLOWLIST | CSRF SYNCHRONISER TOKEN | CSP SELF-ONLY | ZIP-SLIP / BOMB REJECTION | XML ENTITIES OFF | YAML TAGS REJECTED | HASH-CHAINED AUDIT | NO TELEMETRY

NOVELTY & PRIOR ART

NO “FIRST / ONLY” CLAIM

# The algorithms are not new. The integration and the evidence discipline are the contribution.

Feature location, traceability recovery, architecture conformance, and assurance cases are mature fields. Reversa already combines traceability, confidence marking, gap preservation, and a SHA-256 manifest. PreFlight’s difference is domain-specific integration, not a new algorithm.

## CLOSEST PRIOR ART

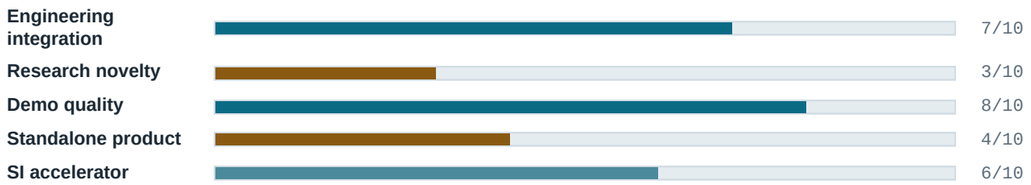
| AREA | REPRESENTATIVE WORK | WHAT IT ESTABLISHES |
|---|---|---|
| **Feature location** | Dit et al. 2013; Marcus et al. 2004 | IR and concept location can recover where a feature lives. |
| **Traceability** | Antoniol et al. 2002; FTLR 2021; LiSSA 2025; TraceLLM 2026 | Links between artefacts can be recovered, ranked, and evaluated. |
| **Reverse documentation** | Reversa, arXiv:2605.18684 (2026) | Traceable specs, confidence marking, gaps for human validation, manifest. |
| **Architecture conformance** | Murphy et al. reflexion models (1995) | Compare intended and observed structure; surface drift. |
| **Assurance cases** | Kelly & Weaver, GSN (2004) | Claims must be supported by argument and evidence. |
| **Tooling** | GoQu, CAST, ADDI, Guidewire CUE | Gosu lint, portfolio discovery, official estimate. |

Selective, not exhaustive. Public IDs and titles only; gated bodies marked gated. Vendor claims labelled as vendor claims. We do not claim the absence of prior work. Citations in `docs/presentation/sources.md`. Reversa is the closest single work; PreFlight is narrower and domain-specific.

**Threat to the novelty claim:** if Reversa adds Guidewire artefact adapters and a disposition ledger, the gap narrows quickly. The moat is not the algorithm. It is the domain evidence model, the review workflow, and the offline independence.

## HONEST SCORECARD

Author’s assessment, not a benchmark or external review.




Research novelty is deliberately low: the individual mechanisms are established. The contribution is the integration, the Guidewire artefact model, and the review artifact.

## WHAT IS ACTUALLY DIFFERENT

1. Cross-artefact feature evidence for Guidewire-style PCF/entity/typelist/template/Gosu.
2. A disposition ledger with per-row provenance and reviewer state.
3. An impact-rated assumption register tied to gaps.
4. A self-verifying review packet, offline, before the official estimate.

VIABILITY

RESEARCH PROTOTYPE

# The wedge is real. The evidence is not yet sufficient to call it a product.

A research prototype today. The defensible path is an independent, offline, pre-CUE evidence packet for one LOB, used as a systems-integrator accelerator — not a replacement for the official estimate or gate.

## KILL RISKS / OPEN EVIDENCE GAPS

- **Synthetic in-sample**
  One 300-file fixture; not a real estate.
- **Internally authored target**
  AI-assisted reference, not blind independent annotation.
- **No human IAA**
  AI-AI concordance would not be human inter-annotator agreement.
- **No real estate**
  Customer code, scale, and naming drift unproven.
- **CUE is official**
  Official Cloud Upgrade Estimate; PreFlight prepares file-level evidence before it.
- **4-feature holdout**
  Naming-shift only; never imported or tuned.
- **Vendor services and assistants**
  Cloud Assurance Services, Lighthouse, and Developer Assistants already own parts of the workflow.
- **Gated rule corpus**
  Public IDs/titles; bodies and threshold are gated.
- **Parser / naming shift**
  Custom PCF shapes and renames can break matching.
- **Draft non-deployable**
  Local _ext draft is not compiled or tenant-validated.

## PROPOSED ACCEPTANCE TESTS BEFORE THE NEXT CLAIM

| TEST | QUESTION | PASS SIGNAL |
|---|---|---|
| Desirability | Would an SI reuse the packet on another LOB? | Requested, not subsidised |
| Feasibility | Do supported parsers cover a real LOB? | ≥ 90% of files parsed |
| Credibility | Do blind humans agree with the links? | IAA κ ≥ 0.6 on 30 links |

## WHAT WE WOULD NEED TO BELIEVE

- An SI will pay for independent pre-estimate evidence, not just delivery hours.
- A real estate produces findings and dispositions comparable to the synthetic slice.
- Reviewers prefer the packet to the spreadsheet they use today.

## DEFENSIBLE WEDGE

### Vendor-independent, offline, pre-CUE, file-level evidence for one LOB.

Why it can survive: independence from the vendor gate, auditability, low deployment friction, and reviewer control. “Independent” means vendor-independent, **not independently validated**: the evaluation is internal and in-sample. The buyer is an SI or an insurer’s migration lead who needs a repeatable evidence pack before CUE and Cloud Assurance Services fix the estimate and the gate.

## VERDICT

**Research prototype today.** Viable as a wedge if a pilot estate and blind human annotation arrive. Not viable as a standalone product on the current evidence.

## ROADMAP TO A CREDIBLE CLAIM

| | | | |
|---|---|---|---|
| **01** Real anonymised estate, one LOB | **02** Second LOB to test transfer | **03** Blind human annotation + IAA | **04** Line-level findings |
| **05** Performance envelope at estate scale | **06** Cloud Assurance / CUE import-export | **07** Compile validation of draft | |

CLOSE

# Every claim has evidence.
# Every unknown has a name.
# Every change waits for a human.

That is the standard PreFlight is trying to make ordinary: an estimate you can inspect, challenge, and sign — with the unknowns written down instead of discovered at week five.

**Pilot estate**
An anonymised, authorised snapshot of one LOB. The next experiment is a real inventory, not another synthetic fixture.

**SME annotation**
Blind review by two domain experts, with disagreements adjudicated. That is the human IAA the current evaluation lacks.

**Feedback**
Attack the claims, the packet, and the disposition model. The fastest way to improve this is to try to break it.

**Pilot ask:** one authorised estate, two blind SME annotators, hard feedback.

**Saundarya Singh** · Independent research prototype **Repository** github.com/alikhan37544/Pre-Flight_Autonomous_Migration **License** MIT **Deck** docs/presentation/index.html

PreFlight is an independent academic research prototype. It is not affiliated with, endorsed by, sponsored by, or certified by Guidewire Software, Inc. Guidewire, InsuranceSuite, PolicyCenter, ClaimCenter, BillingCenter, Gosu, Advanced Product Designer (APD), Jutro, Integration Gateway, Cloud Assurance, Developer Assistants, Guidewire Cloud Platform, and InsuranceNow are trademarks or registered trademarks of Guidewire Software, Inc. All fixtures are synthetic. Findings are pattern-based heuristics, not official Cloud Standards checks, and are not legal, security, or compliance advice. Impact ranges are heuristic planning ranges, not quotes. **Official Cloud Assurance: not run.**